

# Class Notes on GANs Models

Jack Li

October 10, 2024

## Contents

<b>1</b>	<b>Introduction</b>	<b>1</b>
<b>2</b>	<b>Basic Math</b>	<b>1</b>
2.1	Probability Theory . . . . .	1
2.2	Linear Algebra . . . . .	1
2.3	Optimization . . . . .	1
<b>3</b>	<b>PyTorch Basics</b>	<b>1</b>
3.1	Tensors . . . . .	1
<b>4</b>	<b>PyTorch training gradients</b>	<b>2</b>
4.1	Autograd . . . . .	3
4.2	Building Neural Networks . . . . .	4
<b>5</b>	<b>GANs Models</b>	<b>4</b>
5.1	Basic GAN . . . . .	4
5.2	DCGAN . . . . .	4
5.3	WGAN . . . . .	4
5.4	CycleGAN . . . . .	4
<b>6</b>	<b>Conclusion</b>	<b>4</b>

## 1 Introduction

Provide an introduction to GANs and their importance in machine learning.

## 2 Basic Math

### 2.1 Probability Theory

- Definitions of probability, random variables, expectation, etc.

### 2.2 Linear Algebra

- Vectors, matrices, eigenvalues, eigenvectors, etc.

### 2.3 Optimization

- Gradient descent, stochastic gradient descent, etc.

## 3 PyTorch Basics

### 3.1 Tensors

- Definition and operations on tensors.

## 4 PyTorch training gradients

Code 1: PyTorch training gradients

```
1  #SECTION: Gradient computation
2
3  # Step 1: Define a simple model
4  model = nn.Linear(1, 1)
5  optimizer = optim.SGD(model.parameters(), lr=0.01)
6
7  # Dummy input and target
8  input = torch.tensor([[1.0]], requires_grad=True)
9  target = torch.tensor([[2.0]])
10
11 # Step 2: Print the initial parameters
12 print("Initial parameters:")
13 for param in model.parameters():
14     print(param.data)
15
16 # Step 3: Forward pass
17 output = model(input)
18 loss = (output - target).pow(2).mean()
19
20 # Step 4: Zero the gradients
21 optimizer.zero_grad()
22
23 # Step 5: Backward pass
24 loss.backward()
25
26 # Step 6: Update the parameters
27 optimizer.step()
28
29 # Step 7: Print the parameters after the update
30 print("\nParameters after one training step:")
31 for param in model.parameters():
32     print(param.data)
```

### 1. Initialize Parameters:

- Assume initial weights  $w$  and bias  $b$  are both 0.
- Model:  $y = wx + b$

### 2. Forward Pass:

- Compute the output:  $\hat{y} = wx + b$
- Given input  $x = 1.0$  and target  $y = 2.0$ :

$$\hat{y} = 0 \cdot 1.0 + 0 = 0$$

### 3. Compute Loss:

- Loss function: Mean Squared Error (MSE)

$$\text{Loss} = \frac{1}{N} \sum_{i=1}^N (\hat{y}_i - y_i)^2$$

- For our single data point:

$$\text{Loss} = (0 - 2.0)^2 = 4.0$$

### 4. Backward Pass (Gradient Calculation):

- Compute gradients of the loss with respect to  $w$  and  $b$ :

$$\frac{\partial \text{Loss}}{\partial w} = 2(\hat{y} - y)x = 2(0 - 2.0) \cdot 1.0 = -4.0$$

$$\frac{\partial \text{Loss}}{\partial b} = 2(\hat{y} - y) = 2(0 - 2.0) = -4.0$$

### 5. Parameter Update:

- Using Stochastic Gradient Descent (SGD) with learning rate  $\eta = 0.01$ :

$$w_{\text{new}} = w - \eta \frac{\partial \text{Loss}}{\partial w} = 0 - 0.01 \cdot (-4.0) = 0.04$$

$$b_{\text{new}} = b - \eta \frac{\partial \text{Loss}}{\partial b} = 0 - 0.01 \cdot (-4.0) = 0.04$$

### 6. Updated Parameters:

- After one training step, the new parameters are:

$$w = 0.04, \quad b = 0.04$$

**Summary** - Initial parameters:  $w = 0, b = 0$  - After one training step:  $w = 0.04, b = 0.04$

## 4.1 Autograd

- Automatic differentiation in PyTorch.

The active selection `gradient.norm(2, dim = 1)` is a PyTorch operation that computes the L2 norm (Euclidean norm) of the `gradient` tensor along a specified dimension. In this case, the dimension specified is `dim = 1`.

$$\Theta = \underset{\Theta}{\operatorname{argmin}} \frac{1}{B} \sum_{i=1}^B \left[ D(z_i, \Theta) - D(y_i, \Theta) \right] + \lambda \left( \left\| \frac{\partial D(y, \Theta)}{\partial y} \right\| - 1 \right)^2$$

Detailed Explanation:

#### 1. L2 Norm (Euclidean Norm):

- - The L2 norm of a vector is a measure of its magnitude and is calculated as the square root of the sum of the squares of its components. Mathematically, for a vector  $v$ , the L2 norm is given by  $\|v\|_2 = \sqrt{\sum v_i^2}$ .
- - In PyTorch, the `norm` function can compute various types of norms, with the L2 norm being specified by the argument 2.

#### 2. Dimension Specification (`dim = 1`):

- - The `dim` argument specifies the dimension along which the norm is computed. In a multi-dimensional tensor, this allows you to compute norms along specific axes.
- - For example, if `gradient` is a 2D tensor (matrix) with shape `[batchsize, numfeatures]`, setting `dim = 1` means that the norm is computed for each row independently. This results in a tensor of shape `[batchsize]`, where each element is the L2 norm of the corresponding row in the original tensor.

Code 2: PyTorch gradient sampling example

```
1 # Define the sampling function
2 def sample_function(x):
3     return torch.sin(x)
4
5
6 # NOTE: Generate sample points with requires_grad=True, and need requires_grad=True
7 # Generate sample points with requires_grad=True
8 x = torch.tensor(
9     np.linspace(0, 2 * np.pi, 100), dtype=torch.float32, requires_grad=True
10 )
11
12 # Define f by sampling from the sample_function
13 f = sample_function(x)
```

```
14  
15 # Compute the gradient of f with respect to x  
16 grad = torch.autograd.grad(outputs=f, inputs=x, grad_outputs=torch.ones_like(f))  
17  
18 print(f"The gradient of f(x) = sin(x) at x = {x} is {grad[0]}")
```

## 4.2 Building Neural Networks

- Layers, activation functions, loss functions, etc.

## 5 GANs Models

### 5.1 Basic GAN

- Architecture: Generator and Discriminator.
- Loss functions: Minimax game.
- Training process.

### 5.2 DCGAN

- Architecture: Convolutional layers.
- Improvements over basic GAN.
- Training tips.

### 5.3 WGAN

- Wasserstein distance.
- Critic network.
- Gradient penalty.

### 5.4 CycleGAN

- Architecture: Cycle consistency loss.
- Applications: Image-to-image translation.

## 6 Conclusion

Summarize the key points and discuss future directions.