

Class Notes on GANs Models

Jack Li

October 12, 2024

Contents

1	Introduction	1
2	Basic Math	1
2.1	Probability Theory	1
2.2	Linear Algebra	1
2.3	Optimization	1
3	PyTorch Basics	1
3.1	Tensors	1
4	PyTorch training gradients	2
4.1	Autograd	3
4.2	Building Neural Networks	4
5	GANs Models	4
5.1	Basic GAN	4
5.2	DCGAN	4
5.3	WGAN	4
5.4	CycleGAN	4
6	Conclusion	4

1 Introduction

Provide an introduction to GANs and their importance in machine learning.

2 Basic Math

2.1 Probability Theory

- Definitions of probability, random variables, expectation, etc.

2.2 Linear Algebra

- Vectors, matrices, eigenvalues, eigenvectors, etc.

2.3 Optimization

- Gradient descent, stochastic gradient descent, etc.

3 PyTorch Basics

3.1 Tensors

- Definition and operations on tensors.

4 PyTorch training gradients

Code 1: PyTorch training gradients

```
1  #SECTION: Gradient computation
2
3  # Step 1: Define a simple model
4  model = nn.Linear(1, 1)
5  optimizer = optim.SGD(model.parameters(), lr=0.01)
6
7  # Dummy input and target
8  input = torch.tensor([[1.0]], requires_grad=True)
9  target = torch.tensor([[2.0]])
10
11 # Step 2: Print the initial parameters
12 print("Initial parameters:")
13 for param in model.parameters():
14     print(param.data)
15
16 # Step 3: Forward pass
17 output = model(input)
18 loss = (output - target).pow(2).mean()
19
20 # Step 4: Zero the gradients
21 optimizer.zero_grad()
22
23 # Step 5: Backward pass
24 loss.backward()
25
26 # Step 6: Update the parameters
27 optimizer.step()
28
29 # Step 7: Print the parameters after the update
30 print("\nParameters after one training step:")
31 for param in model.parameters():
32     print(param.data)
```

1. Initialize Parameters:

- Assume initial weights w and bias b are both 0.
- Model: $y = wx + b$

2. Forward Pass:

- Compute the output: $\hat{y} = wx + b$
- Given input $x = 1.0$ and target $y = 2.0$:

$$\hat{y} = 0 \cdot 1.0 + 0 = 0$$

3. Compute Loss:

- Loss function: Mean Squared Error (MSE)

$$\text{Loss} = \frac{1}{N} \sum_{i=1}^N (\hat{y}_i - y_i)^2$$

- For our single data point:

$$\text{Loss} = (0 - 2.0)^2 = 4.0$$

4. Backward Pass (Gradient Calculation):

- Compute gradients of the loss with respect to w and b :

$$\frac{\partial \text{Loss}}{\partial w} = 2(\hat{y} - y)x = 2(0 - 2.0) \cdot 1.0 = -4.0$$

$$\frac{\partial \text{Loss}}{\partial b} = 2(\hat{y} - y) = 2(0 - 2.0) = -4.0$$

5. Parameter Update:

- Using Stochastic Gradient Descent (SGD) with learning rate $\eta = 0.01$:

$$w_{\text{new}} = w - \eta \frac{\partial \text{Loss}}{\partial w} = 0 - 0.01 \cdot (-4.0) = 0.04$$

$$b_{\text{new}} = b - \eta \frac{\partial \text{Loss}}{\partial b} = 0 - 0.01 \cdot (-4.0) = 0.04$$

6. Updated Parameters:

- After one training step, the new parameters are:

$$w = 0.04, \quad b = 0.04$$

Summary - Initial parameters: $w = 0, b = 0$ - After one training step: $w = 0.04, b = 0.04$

4.1 Autograd

- Automatic differentiation in PyTorch.

The active selection `gradient.norm(2, dim = 1)` is a PyTorch operation that computes the L2 norm (Euclidean norm) of the `gradient` tensor along a specified dimension. In this case, the dimension specified is `dim = 1`.

$$\Theta = \underset{\Theta}{\operatorname{argmin}} \frac{1}{B} \sum_{i=1}^B \left[D(z_i, \Theta) - D(y_i, \Theta) \right] + \lambda \left(\left\| \frac{\partial D(y, \Theta)}{\partial y} \right\| - 1 \right)^2$$

Detailed Explanation:

1. L2 Norm (Euclidean Norm):

- - The L2 norm of a vector is a measure of its magnitude and is calculated as the square root of the sum of the squares of its components. Mathematically, for a vector v , the L2 norm is given by $\|v\|_2 = \sqrt{\sum v_i^2}$.
- - In PyTorch, the `norm` function can compute various types of norms, with the L2 norm being specified by the argument 2.

2. Dimension Specification (`dim = 1`):

- - The `dim` argument specifies the dimension along which the norm is computed. In a multi-dimensional tensor, this allows you to compute norms along specific axes.
- - For example, if `gradient` is a 2D tensor (matrix) with shape `[batchsize, numfeatures]`, setting `dim = 1` means that the norm is computed for each row independently. This results in a tensor of shape `[batchsize]`, where each element is the L2 norm of the corresponding row in the original tensor.

Code 2: PyTorch gradient sampling example

```
1 # Define the sampling function
2 def sample_function(x):
3     return torch.sin(x)
4
5
6 # NOTE: Generate sample points with requires_grad=True, and need requires_grad=True
7 # Generate sample points with requires_grad=True
8 x = torch.tensor(
9     np.linspace(0, 2 * np.pi, 100), dtype=torch.float32, requires_grad=True
10 )
11
12 # Define f by sampling from the sample_function
13 f = sample_function(x)
```

```

14
15 # Compute the gradient of f with respect to x
16 grad = torch.autograd.grad(outputs=f, inputs=x, grad_outputs=torch.ones_like(f))
17
18 print(f"The gradient of f(x) = sin(x) at x = {x} is {grad[0]}")

```

4.2 Building Neural Networks

- Layers, activation functions, loss functions, etc.

4.3 Loss Functions

— PyTorch Conv2d Equation

The output size of a Conv2d layer can be calculated using the following equation:

$$\text{Output Size} = \left\lfloor \frac{\text{Input Size} + 2 \times \text{Padding} - \text{Kernel Size}}{\text{Stride}} \right\rfloor + 1$$

Where: - Input Size is the size of the input feature map (height or width). - Padding is the number of zero-padding added to both sides of the input. - Kernel Size is the size of the convolution kernel (height or width). - Stride is the stride of the convolution.

PyTorch ConvTranspose2d Equation

The output size of a ConvTranspose2d (transposed convolution) layer can be calculated using the following equation:

$$\text{Output Size} = (\text{Input Size} - 1) \times \text{Stride} - 2 \times \text{Padding} + \text{Kernel Size} + \text{Output Padding}$$

Where: - Input Size is the size of the input feature map (height or width). - Stride is the stride of the convolution. - Padding is the number of zero-padding added to both sides of the input. - Kernel Size is the size of the convolution kernel (height or width). - Output Padding is the additional size added to the output (usually used to ensure the output size matches a specific value).

5 GANs Models

5.1 Basic GAN

- Architecture: Generator and Discriminator.
- Loss functions: Minimax game.
- Training process.

5.2 DCGAN

- Architecture: Convolutional layers.
- Improvements over basic GAN.
- Training tips.

5.3 WGAN

- Wasserstein distance.
- Critic network.
- Gradient penalty.

5.4 CycleGAN

- Architecture: Cycle consistency loss.
- Applications: Image-to-image translation.

6 Conclusion

Summarize the key points and discuss future directions.